

Red Flag Use Case Triage

Purpose: Sort an AI workflow into a practical green, yellow, or red category before launch, sale, demo, automation, or client delivery.

When To Use It

Use this before launching, selling, demoing, advertising, automating, or delivering an AI workflow that could affect people, sensitive data, professional advice, client obligations, or real-world outcomes.

This worksheet does not decide whether a use is allowed. It helps you spot uses that deserve more caution, narrower scope, human review, or qualified help.

What This Produces

You should finish with:

- a short workflow description
- a green/yellow/red triage decision
- red-flag categories marked Yes, No, or Unsure
- follow-up actions for anything that needs review

Workflow Description

Question	Answer
What are you using AI to do?	
Who sees or relies on the output?	[Only me / collaborators / clients / customers / users / subscribers / public / other]
Does a person review the output before it is used?	[Yes / No / Sometimes / Unsure]
Does the workflow use customer, client, sensitive, confidential, or restricted data?	[Yes / No / Unsure]
Is the workflow public, paid, advertised, or client-facing?	[Yes / No / Planned / Unsure]

Green / Yellow / Red Sort

Green: basic records are probably enough for now

Use green when the workflow is low stakes, uses no sensitive data, and a human reviews the result.

Examples:

- brainstorming ideas
- outlining posts, products, or courses
- editing your own draft
- summarizing public sources for internal use
- creating internal checklists
- generating placeholder copy that will be rewritten
- coding help on non-sensitive personal projects

Required records:

- ☐ tool listed in project snapshot or vendor note
- ☐ no sensitive, client, confidential, credential, or restricted data used
- ☐ human review before publishing, delivery, or customer use

Yellow: review before selling, publishing, or delivering

Use yellow when a buyer, client, customer, user, or audience may rely on the output, or when platform, client, data, disclosure, or claims questions may matter.

Examples:

- client deliverables
- customer-facing copy
- paid templates, guides, prompts, courses, downloads, or tools
- AI-generated images, audio, video, or synthetic media
- advice-like content
- code used in a paid product or client project
- AI support reply drafts
- AI-assisted research summaries
- real names, brands, likenesses, voices, testimonials, reviews, or protected works

Required records:

- ☐ claims checked
- ☐ disclosure decision recorded
- ☐ data rules checked
- ☐ platform/client terms checked where relevant
- ☐ human review completed or clearly assigned to the user
- ☐ open questions listed

Red: pause, narrow, or get qualified review

Use red when the workflow could affect important rights, sensitive data, regulated decisions, professional duties, safety, or access to opportunities.

Examples:

- legal, financial, medical, tax, mental-health, employment, education, housing, insurance, safety, or similar advice
- hiring, scoring, eligibility, pricing, lending, insurance, school, benefits, fraud, identity, or access decisions
- sensitive customer, client, health, financial, legal, child, student, employee, biometric, credential, or confidential data
- AI impersonation, voice cloning, likeness use, or deepfakes involving real people
- scraping, copying, or training on content where rights or terms are unclear
- promises of legal compliance, guaranteed outcomes, expert review, or professional certainty
- work restricted by client contract, NDA, professional rule, platform policy, or privacy promise

Required action:

- ☐ stop, pause, or narrow the workflow
- ☐ identify the specific question
- ☐ ask a qualified reviewer, platform, client, vendor, or other appropriate source
- ☐ save the answer before proceeding

Red-Flag Categories

Category	Applies?	What to check	Next step
Health or medical	[Yes / No / Unsure]	Could users treat output as health guidance? Is health data processed?	Pause or narrow; get qualified health/privacy review.
Mental health or crisis support	[Yes / No / Unsure]	Could users seek therapy, diagnosis, crisis help, or emotional support?	Pause high-risk features; route users to appropriate help where relevant.
Legal	[Yes / No / Unsure]	Could output be mistaken for legal advice or document-specific guidance?	Reframe as general drafting/support or route to counsel.
Financial, tax, credit, or insurance	[Yes / No / Unsure]	Does AI recommend money actions, pricing, eligibility, loans, tax treatment, claims, or coverage?	Narrow scope; get qualified review if decisions or advice are involved.
Employment, HR, hiring, education, or housing	[Yes / No / Unsure]	Does AI rank, screen, grade, evaluate, approve, deny, or recommend outcomes?	Pause automated decisions; get HR/legal/sector review.
Children or minors	[Yes / No / Unsure]	Could children use it or appear in the data?	Narrow access and get privacy/legal review.
Biometrics, identity, surveillance, or fraud	[Yes / No / Unsure]	Does AI process face, voice, identity, monitoring, fraud, or access signals?	Pause and get privacy/legal/security review.
Sensitive personal data	[Yes / No / Unsure]	What sensitive data enters prompts, files, logs, analytics, or outputs?	Minimize data; avoid unnecessary processing; review data handling.
Client confidential or contract-restricted data	[Yes / No / Unsure]	Do contracts, NDAs, client policies, or platform rules restrict AI use?	Check permissions before using AI on the data.
Scraping, third-party data, IP, or likeness	[Yes / No / Unsure]	Are source rights, platform terms, consent, or reuse rights unclear?	Review source rights, terms, and consent before launch.
Automated decisions with real-world consequences	[Yes / No / Unsure]	Does AI decide, or only assist? Can a person override or appeal?	Keep AI advisory; add human review; document decision path.

Triage Decision

Decision	Mark one	Notes
Green: proceed with basic records	[]	
Yellow: proceed only after listed review steps	[]	
Red: pause, narrow, or get qualified help	[]	

Reason:

[NOTES]

Follow-Up Actions

Action	Owner	Due date	Status
			[Open / In progress / Done / Blocked]
			[Open / In progress / Done / Blocked]

Review Carefully If

- Any category is marked Yes or Unsure.
- The AI could affect a person's rights, money, health, safety, job, education, housing, insurance, credit, identity, or access to services.
- You are using data from children, sensitive personal data, biometric data, scraped data, or client confidential information.
- Users may think the AI has authority, expertise, or final decision-making power.

Short boundary: This worksheet is educational information, not legal advice. It helps you spot red flags and decide what to review before launch. Review specific legal, privacy, contractual, regulatory, or professional obligations with qualified counsel or the appropriate advisor.