

AI Change Log

Purpose: Track changes to the model, prompt, vendor, data use, claims, disclosures, support process, and product behavior over time.

When To Use It

Use this whenever you change anything that could affect what the AI does, what data it touches, what users see, what you claim, what support needs to know, or what records you rely on.

What This Produces

You should finish with a simple running record of meaningful changes, starting with the launch baseline.

Add one row per meaningful change. Small typo fixes may not need a row. Anything that affects output behavior, data flow, vendor dependency, public copy, disclosure, FAQ, human review, customer support, or review scope should be logged.

Change Types

- Launch baseline
- Model change
- Vendor change
- Prompt change
- System instruction change
- Data-use change
- Logging or analytics change
- Feature behavior change
- Public claim change
- Disclosure or FAQ change
- Human review change
- Red-flag scope change
- Incident-driven change
- Support process change
- Other

Change Log

Date	Change type	What changed	Why it changed	Risk impact	User-facing impact	Docs updated	Tested or checked by	Next review
				[Lower / Same / Higher / Unknown]		[Snapshot / Claims / Data intake / Disclosure / FAQ / Vendor note / Incident log / Other]		
				[Lower / Same / Higher / Unknown]				
				[Lower / Same /				

				Higher / Unknown]				
--	--	--	--	-------------------	--	--	--	--

Launch Baseline Entry

Create a first entry when the project launches or enters beta.

Field	Answer
Launch or beta date	
Initial model/vendor	
Initial prompt or instruction location	
Initial data flow	
Initial claims reviewed	
Initial disclosure location	
Initial customer FAQ location	
Initial human review approach	
Known limitations at launch	
First review date	

Review Prompts Before A Change Goes Live

- Does this change what users should be told?
- Does this change what data goes into the AI workflow?
- Does this change prompt logs, analytics, retention, or vendor exposure?
- Does this change the accuracy, tone, risk, or reliability of output?
- Does this change whether human review is needed?
- Does this make any public claim outdated?
- Does this create a new red-flag category?
- Does this require a customer FAQ, disclosure, or support-script update?
- Does this require qualified review?

Review Carefully If

- You changed the vendor, model, prompt, system instruction, data source, logging behavior, or automation path.
- You changed public claims, pricing pages, demos, onboarding copy, proposals, or marketplace listings.
- You removed human review or made output more automated.
- The change was made because of a complaint, bad output, privacy concern, platform notice, or customer confusion.

Short boundary: This worksheet is educational information, not legal advice. It helps you keep AI changes visible and spot questions to review. Review specific legal, privacy, contractual, security, platform, or regulatory obligations with qualified counsel or the appropriate advisor.